Study Guide: Data Retrieval with SQL

Afshine Amidi and Shervine Amidi

August 21, 2020

General concepts

❏ **Structured Query Language** – Structured Query Language, abbreviated as SQL, is a language that is largely used in the industry to query data from databases.

❏ **Query structure** – Queries are usually structured as follows:

```SQL
-- Select fields              mandatory
SELECT
    col_1,
    col_2,
    ... ,
    col_n

-- Source of data            mandatory
FROM table t

-- Gather info from other sources    optional
JOIN other_table ot
  ON (t.key = ot.key)

-- Conditions               optional
WHERE some_condition(s)

-- Aggregating              optional
GROUP BY column_group_list

-- Sorting values           optional
ORDER BY column_order_list

-- Restricting aggregated values   optional
HAVING some_condition(s)

-- Limiting number of rows     optional
LIMIT some_value
```

Remark: the SELECT DISTINCT *command can be used to ensure not having duplicate rows.*

❏ **Condition** – A condition is of the following format:

```SQL
some_col some_operator some_col_or_value
```

where some_operator can be among the following common operations:

Category	Operator	Command
General	Equality / non-equality	= / !=, <>
	Inequalities	>=, >, <, <=
	Belonging	IN (val_1, ..., val_n)
	And / or	AND / OR
	Check for missing value	IS NULL
	Between bounds	BETWEEN val_1 AND val_2
Strings	Pattern matching	LIKE '%val%'

❏ **Joins** – Two tables table_1 and table_2 can be joined in the following way:

```SQL
...

FROM table_1 t1
type_of_join table_2 t2
  ON (t2.key = t1.key)

...
```

where the different type_of_join commands are summarized in the table below:

Type of join	Illustration
INNER JOIN	
LEFT JOIN	
RIGHT JOIN	
FULL JOIN	

Remark: joining every row of table 1 with every row of table 2 can be done with the CROSS JOIN *command, and is commonly known as the cartesian product.*

Aggregations

❏ **Grouping data** – Aggregate metrics are computed on grouped data in the following way:

The SQL command is as follows:

```
SQL
SELECT
    col_1,
    agg_function(col_2)
FROM table
GROUP BY col_1
```

❏ **Grouping sets** – The GROUPING SETS command is useful when there is a need to compute aggregations across different dimensions at a time. Below is an example of how all aggregations across two dimensions are computed:

```
SQL
SELECT
    col_1,
    col_2,
    agg_function(col_3)
FROM table
GROUP BY (
  GROUPING SETS
    (col_1),
    (col_2),
    (col_1, col_2)
)
```

❏ **Aggregation functions** – The table below summarizes the main aggregate functions that can be used in an aggregation query:

Category	Operation	Command
Values	Mean	AVG(col)
	Percentile	PERCENTILE_APPROX(col, p)
	Sum / # of instances	SUM(col) / COUNT(col)
	Max / min	MAX(col) / MIN(col)
	Variance / standard deviation	VAR(col) / STDEV(col)
Arrays	Concatenate into array	collect_list(col)

*Remark: the median can be computed using the PERCENTILE_APPROX function with p equal to 0.5.*

❏ **Filtering** – The table below highlights the differences between the WHERE and HAVING commands:

WHERE	HAVING
- Filter condition applies to individual rows	- Filter condition applies to aggregates
- Statement placed right after FROM	- Statement placed right after GROUP BY

Remark: if WHERE and HAVING are both in the same query, WHERE will be executed first.

Window functions

❏ **Definition** – A window function computes a metric over groups and has the following structure:

The SQL command is as follows:

```
SQL
some_window_function() OVER(PARTITION BY some_col ORDER BY another_col)
```

Remark: window functions are only allowed in the SELECT clause.

❏ **Row numbering** – The table below summarizes the main commands that rank each row across specified groups, ordered by a specific column:

Command	Description	Example
ROW_NUMBER()	Ties are given different ranks	1, 2, 3, 4
RANK()	Ties are given same rank and skip numbers	1, 2, 2, 4
DENSE_RANK()	Ties are given same rank and don't skip numbers	1, 2, 2, 3

❏ **Values** – The following window functions allow to keep track of specific types of values with respect to the partition:

Command	Description
FIRST_VALUE(col)	Takes the first value of the column
LAST_VALUE(col)	Takes the last value of the column
LAG(col, n)	Takes the n^{th} previous value of the column
LEAD(col, n)	Takes the n^{th} following value of the column
NTH_VALUE(col, n)	Takes the n^{th} value of the column

Advanced functions

❒ **SQL tips** – In order to keep the query in a clear and concise format, the following tricks are often done:

Operation	Command	Description
Renaming columns	`SELECT operation_on_column AS col_name`	New column names shown in query results
Abbreviating tables	`FROM table_1 t1`	Abbreviation used within query for simplicity in notations
Simplifying group by	`GROUP BY col_number_list`	Specify column position in `SELECT` clause instead of whole column names
Limiting results	`LIMIT n`	Display only n rows

❒ **Sorting values** – The query results can be sorted along a given set of columns using the following command:

SQL
```
... [query] ...
ORDER BY col_list
```

Remark: by default, the command sorts in ascending order. If we want to sort it in descending order, the `DESC` command needs to be used after the column.

❒ **Column types** – In order to ensure that a column or value is of one specific data type, the following command is used:

SQL
```
CAST(some_col_or_value AS data_type)
```

where `data_type` is one of the following:

Data type	Description	Example
`INT`	Integer	`2`
`DOUBLE`	Numerical value	`2.0`
`STRING`	String	`'teddy bear'`
`VARCHAR`		
`DATE`	Date	`'2020-01-01'`
`TIMESTAMP`	Timestamp	`'2020-01-01 00:00:00.000'`

Remark: if the column contains data of different types, the `TRY_CAST()` command will convert unknown types to `NULL` instead of throwing an error.

❒ **Column manipulation** – The main functions used to manipulate columns are described in the table below:

Category	Operation	Command
General	Take first non-`NULL` value	`COALESCE(col_1, col_2, ..., col_n)`
	Create a new column combining existing ones	`CONCAT(col_1, ..., col_n)`
Value	Round value to n decimals	`ROUND(col, n)`
String	Converts string column to lower / upper case	`LOWER(col) / UPPER(col)`
	Replace occurrences of old in col to new	`REPLACE(col, old, new)`
	Take the substring of col, with a given start and length	`SUBSTR(col, start, length)`
	Remove spaces from the left / right / both sides	`LTRIM(col) / RTRIM(col) / TRIM(col)`
	Length of the string	`LENGTH(col)`
Date	Truncate at a given granularity (year, month, week)	`DATE_TRUNC(time_dimension, col_date)`
	Transform date	`DATE_ADD(col_date, number_of_days)`

❒ **Conditional column** – A column can take different values with respect to a particular set of conditions with the `CASE WHEN` command as follows:

SQL
```
CASE WHEN some_condition THEN some_value
         ...
     WHEN some_other_condition THEN some_other_value
     ELSE some_other_value_n END
```

❒ **Combining results** – The table below summarizes the main ways to combine results in queries:

Category	Command	Remarks
Union	`UNION`	Guarantees distinct rows
	`UNION ALL`	Potential newly-formed duplicates are kept
Intersection	`INTERSECT`	Keeps observations that are in all selected queries

❒ **Common table expression** – A common way of handling complex queries is to have temporary result sets coming from intermediary queries, which are called common table expressions (abbreviated CTE), that increase the readability of the overall query. It is done thanks to the `WITH ... AS ...` command as follows:

SQL
```
WITH cte_1 AS (
SELECT ...
),
```

```
...

cte_n AS (
SELECT ...
)

SELECT ...
FROM ...
```

Command	Description
OVERWRITE	Overwrites existing data
INTO	Appends to existing data

❒ **Dropping table** – Tables are dropped in the following way:

SQL
```
DROP TABLE table_name;
```

Table manipulation

❒ **Table creation** – The creation of a table is done as follows:

SQL
```
CREATE [table_type] TABLE [creation_type] table_name(
  col_1 data_type_1,
       ...        ,
  col_n data_type_n
)
[options];
```

❒ **View** – Instead of using a complicated query, the latter can be saved as a view which can then be used to get the data. A view is created with the following command:

SQL
```
CREATE VIEW view_name AS complicated_query;
```

Remark: a view does not create any physical table and is instead seen as a shortcut.

where [table_type], [creation_type] and [options] are one of the following:

Category	Command	Description
Table type	Blank	Default table
	EXTERNAL TABLE	External table
Creation type	Blank	Creates table and overwrites current one if it exists
	IF NOT EXISTS	Only creates table if it does not exist
Options	location 'path_to_hdfs_folder'	Populate table with data from hdfs folder
	stored as data_format	Stores the table in a specific data format, e.g. parquet, orc or avro

❒ **Data insertion** – New data can either append or overwrite already existing data in a given table as follows:

SQL
```
WITH ...                        -- optional

INSERT [insert_type] table_name    -- mandatory

SELECT ...;                     -- mandatory
```

where [insert_type] is among the following: